
Experience Admission for White-Box Agentic RL: A Position on Dataflow between Exploration and the Base Policy

Anonymous Author(s)
Affiliation
Address
email

Abstract

1 In white-box instrumented-worker multi-worker RL, a layer is often miss-
2 ing between optimizers and collaboration surfaces: after trajectories are
3 produced, who may materialize, at what depth, and when may θ update.
4 Black-box APIs are out of strong-form claim.
5 We call this layer experience admission. Novelty is P1–P4 (a proposed
6 admission policy) plus preregistered failure conditions—not a new score
7 or proven discovery. Hypothesis-level claim: the bottleneck is often how
8 experience enters the base update path in a layered, isolated, low-frequency
9 way.
10 Layered test: fix $\mathcal{A}$; vary only the data plane (A/B/P). If failures 4–5 do
11 not fire, materialization cost falls vs. A, and the task is not in persistent
12 regression, meet the scale-up gate only (allows larger runs; $\neq$ validation;
13 does not alone entail “must ship as a first-class module”). After scale-up
14 still passes, then recommend a first-class pipeline module; else withdraw
15 that configuration’s strong form.

16 1 Motivation (scoped)

17 As agent environments scale—more workers, longer horizons, denser trajectory streams—
18 three layers are often mixed: optimizer (when/how to update; e.g., group-relative meth-
19 ods [1]; single-rollout async optimization [2]), collaboration (shared context at interaction
20 time), and admission (who may materialize, how deeply, when θ may be touched). Stability
21 $\neq$ admission: the former asks whether updates diverge on a visible set; the latter asks which
22 trajectories may enter that set and whether/how they are materialized first. IMPALA [3]
23 can still treat “delivered $\Rightarrow$ learnable.” Shielding [4] constrains actions; CQL [5] constrains
24 offline values—not online materialization/update budgets. Scaling the environment with-
25 out an admission layer can amplify unbudgeted exposure of noise and hazard into the base
26 policy.

27 Minimal counterexample. With ordinary learnable τ^+ and high-TD-error hazard $\tau^\dagger$ in one
28 buffer, PER [6] allows $\tau^\dagger$ into the learner above uniform rate when high-priority (power-law
29 sampling: high quantile does not guarantee a draw, but the mechanism does not exclude
30 hazard uplift). Action shields do not remove it from the train-visible set; delete filters discard
31 analysis value. Gap triad: P1 budget semantics absent (pool-default learnable; no default-
32 off materialization), plus P2-P4 / P3 violations. If hazard has low TD-error, PER may
33 ignore it—that is “unnoticed,” not policy-level quarantine-with-retention (P2). Sampling
34 weights $\neq$ admission policy.

Table 1: Proposed constraints P1–P4 vs. typical PER / delete-style filters.

	Meaning	vs PER / delete-filter
P1	Full materialization off by default (explicit quota)	Usually absent
P2	Quarantine $\neq$ delete	Usually absent
P3	Full materialization $\neq$ admission	Sampled $\Rightarrow$ learned
P4	$D_{\text{adm}} \subsetneq D_{\text{all}}$ when hot nonempty	Reweight-all / action-only

35 Scope. Strong form assumes workers expose hidden representations, local action distribu-
 36 tions, or uncertainty signals. Fingerprints and thresholds are knobs, not results. Black-box
 37 APIs degrade to post-hoc text filtering and must not be conflated with the strong form.

38 2 Proposed admission policy (not “irreducible discovery”)

39 Terms. Materialization = persist a trajectory from a streaming state into an inspectable
 40 object; deep read = costly analysis on (usually) materialized trajectories. They often share
 41 an admission-cost budget but are logically separable. For brevity we say materialization
 42 cost in budget/metrics. D_{read} is eligible for materialize/deep-read, not admit-to-update.

43 Use the name only if all four are enforced; else revert terms. P1–P4 are a proposed admission
 44 policy, not an established scientific finding.

45 Logical interdependence (design-level, not ablation). Drop P1 $\rightarrow$ materialization budget
 46 vanishes; warm/hot become post-hoc labels $\approx$ unbudgeted PER + tags. Drop P2 $\rightarrow$ delete-
 47 or-learn (no analysis track). Drop P3 $\rightarrow D_{\text{read}} \equiv D_{\text{adm}} \approx$ PER/passive pool. Drop P4 $\rightarrow$
 48 P1–P3 may remain, but when hot is nonempty $D_{\text{adm}} \subsetneq D_{\text{all}}$ is no longer forced: a hybrid
 49 cache still distinct from PER, yet incomplete as a visible-set policy. Together they block
 50 four distinct collapses; joint proposal is warranted. This does not replace causal ablations
 51 (§5).

52 LLM curation (filtering / RLAIF / RAFT [7–9]) optimizes pre-batch quality; kept $\Rightarrow$ update
 53 θ . High reward $\neq$ low risk. Multi-phase exploration (e.g., First Return, Then Explore [10])
 54 targets exploration schedules and coverage—orthogonal to materialization budgets, quaran-
 55 tine tracks, or enforced $D_{\text{adm}} \subsetneq D_{\text{all}}$.

56 Honest empirical gap: no causal ablation of which constraint drives gains vs. PER. H3
 57 (measurable quarantine utility) is exploratory for the scale-up gate; P2 execution integrity
 58 (nonempty quarantine must be read/used) is enforced by failure 4(b).

59 $H \leftrightarrow P$: H1 $\rightarrow$ P1 (primary); H2 $\rightarrow$ P2+P4 (primary); H3 $\rightarrow$ P2 value magnitude (exploratory);
 60 H4 (contact frequency / visible-set size, not SGD step size) $\rightarrow$ P3+P4 (primary).

61 3 Interface, predicate, contract

62 Not a full formal system—an interface plus minimal predicates. Dataflow (pool $\mathcal{P}$, not batch
 63 B_t):

$$\tau \xrightarrow{\phi} x \xrightarrow{r} \mathcal{P} \xrightarrow{\text{adm}} D_{\text{adm}} \xrightarrow{\mathcal{A}} \theta.$$

64

$$\tau \in D_{\text{adm}} \iff \text{Validated}(\tau) \wedge \neg \text{HotHazard}(\tau) \wedge \text{InBudget}(\tau),$$

65 with $\text{HotHazard} \iff (r=\text{hot} \wedge \ell=\text{hazard})$. Operational floor for Validated: promotion
 66 must pass a check independent of the routing score that proposed promotion (held-out
 67 probe, second scorer, or human/audit gate—fixed in the experimental plan). Hazard labels
 68 must not be dictated solely by the same routing score. Interaction: every optimizer batch
 69 $B_t \subseteq D_{\text{adm}}$. $B_t \not\subseteq D_{\text{adm}}$ is a contract breach; if the batch contains hot-hazard, failure 4 also
 70 fires.

71 Contract: (1) no hot-hazard in update batches; (2) materialization $\neq$ write θ ; (3) promo-
 72 tion validated independently of routing score; (4) archive closed by default; warm bud-
 73 geted; (5) optional curator cannot enlarge allowed sets. Strong form forbids admit-without-
 74 materialization as default.

75 4 Preregistered falsification

76 Arms: A passive (default materializable / replay); B admission (P1–P4+§3); P optional
 77 PER (weights $\neq$ admission).

78 Metrics: materialization ratio; wall-clock/FLOPs; return; hazard penetration; materializa-
 79 tion gain = probe return lift or loss drop after materialization (pick one; fixed in plan)—
 80 subset mean/median or preregistered smoothing allowed; Spearman ρ vs. (aggregated) gain
 81 for failure-5 labels only.

82 Failures ($K \geq 3$ segments; any one rejects the setting):

- 83 1. Warm-mass collapse—warm mass (routing count) $\geq 90\%$ (or archive+hot $< 10\%$) and
 84 materialization improvement vs. A $< 5\%$ (conjunction; high mass with clear cost drop
 85 does not fire).
- 86 2. No budget gain—materialization improvement $< 10\%$ and wall-clock/FLOPs improve-
 87 ment $< 5\%$ (default AND). Grey zone: mat. gain in $[5\%, 10\%)$ with wall-clock $\geq 5\%$
 88 fires neither 1–2 nor meets the scale-up gate—no scale-up claim; alternative joint criteria
 89 require preregistration.
- 90 3. Persistent regression vs. A (preregistered α /MID).
- 91 4. Hard (P2 integrity): (a) any hot-hazard enters D_{adm} and updates θ ; or (b) quarantine/hot
 92 nonempty in-window but never read/deep-read and unused for any preregistered periph-
 93 eral analysis (“quarantine-to-forget”). (b) N/A if hot empty.
- 94 5. Hard proxy-gain test (three clauses; current proxy family only; any clause firing fails—
 95 run all three):
 96 • Correlation: one-sided Spearman of $\rho > 0$ at preregistered α (e.g. 0.05); nonsignificant
 97 $\rightarrow$ fires.
 98 • Random baseline: archive F1/accuracy within preregistered δ of a matched-marginal
 99 i.i.d. random archive labeling $\rightarrow$ fires.
 100 • Family boundary: “new family” = architectural score change (e.g. entropy $\rightarrow$ prediction
 101 error); retuning thresholds/windows inside the same function = same instance (must
 102 retest). After fail: at most weak post-hoc filtering; no strong-form mix; new family
 103 requires fresh preregistration.

104 Scale-up gate ($\neq$ validation): materialization $\downarrow \geq 10\%$ vs. A; failure 3 unmet; failures 4–
 105 5 unmet (incl. 4(b)). Allows larger experiments only—does not alone entail a first-class
 106 module. After scale-up still passes: then recommend first-class pipeline module. H3 value
 107 magnitude remains exploratory.

108 5 Limitations

109 No theory; no causal minimality claim; no black-box strong form; no LoRA/GRPO main
 110 results; no appendix-as-proof. Condition 5 is intentionally strict: under small N or high
 111 label noise, nonsignificance is likely and should be read as “this proxy family is unsup-
 112 ported,” not “the admission question is void.” Materialization-gain labels are measured on
 113 trajectories selected for materialization, so selection bias is intrinsic; the random-baseline
 114 clause partially disciplines it but does not remove it. Scope boundary: long-horizon, costly
 115 materialization/deep-read, heterogeneous workers; short trajectories / single worker / negli-
 116 gible cost may favor passive pool or PER—even after a pass, designers may decline to ship.
 117 Future: ablations; open-weight multi-seed; learnable routers; H3 value metrics beyond 4(b);
 118 separate paper for black-box weak form [11–13].

6 Conclusion

Deliverables: name the (white-box) layer; propose P1–P4 with logical interdependence; give a falsifiable protocol. Scale-up gate $\neq$ validation and $\neq$ must-ship; after scale-up still passes, then recommend a first-class pipeline module. Whether the hypothesis holds at scale remains open—outside any “already proven” claim.

References

- [1] Z. Shao, P. Wang, Q. Zhu, et al. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. arXiv preprint arXiv:2402.03300, 2024.
- [2] Z. Hou, Y. Li, J. Tang, and Y. Dong. Single-rollout asynchronous optimization for agentic reinforcement learning. arXiv preprint arXiv:2607.07508, 2026.
- [3] L. Espeholt et al. IMPALA: Scalable distributed deep-RL with importance weighted actor-learner architectures. In International Conference on Machine Learning (ICML), 2018.
- [4] M. Alshiekh et al. Safe reinforcement learning via shielding. In Proceedings of the AAAI Conference on Artificial Intelligence, 2018.
- [5] A. Kumar, A. Zhou, G. Tucker, and S. Levine. Conservative Q-learning for offline reinforcement learning. In Advances in Neural Information Processing Systems (NeurIPS), 2020.
- [6] T. Schaul, J. Quan, I. Antonoglou, and D. Silver. Prioritized experience replay. In International Conference on Learning Representations (ICLR), 2016.
- [7] Y. Hu et al. Towards comprehensive preference data collection for reward modeling. arXiv preprint arXiv:2406.16486, 2024.
- [8] H. Lee et al. RLAIIF vs. RLHF. arXiv preprint arXiv:2309.00267, 2023.
- [9] H. Dong et al. RAFT: Reward ranked finetuning. Transactions on Machine Learning Research, 2023.
- [10] A. Ecoffet et al. First return, then explore. Nature, 590:580–586, 2021.
- [11] T. Shen et al. A comprehensive survey of LLM alignment techniques. arXiv preprint arXiv:2407.16216, 2024.
- [12] D. Pathak, P. Agrawal, A. Efros, and T. Darrell. Curiosity-driven exploration by self-supervised prediction. In International Conference on Machine Learning (ICML), 2017.
- [13] Y. Burda, H. Edwards, A. Storkey, and A. Klimov. Exploration by random network distillation. In International Conference on Learning Representations (ICLR), 2019.